

# Lecture 21-22

Hash Tables

# Hash Tables

# Recap of BSTs

- Balanced BSTs offer  $O(\log n)$  time costs for add/remove/find operations by exploiting order relationships among data.
- They are also memory efficient, in that only as many nodes are allocated as elements are contained in the BST
- However...

...by utilizing more memory (greater space complexity), we can achieve even **higher** performance (lower time complexity)

**Hash sets, hash tables and hash maps** offer  **$O(1+\alpha^*)$**  time costs for **add/remove/find** operations by investing more memory in the underlying storage.

\* load factor

# DSC30 gang

- (Marina, 1)  $\leftarrow$  The Boss
- (Kalkin, 2)  $\leftarrow$  Tutor 1
- (Cassidy, 3)  $\leftarrow$  Tutor 2
- .....
- (Elif, 25)  $\leftarrow$  Student 1
- (Grace, 26)  $\leftarrow$  Student 2
- .....

How to store them?

# Direct hashing

- (Marina, 1)  $\leftarrow$  The Boss
- (Kalkin, 2)  $\leftarrow$  Tutor 1
- (Cassidy, 3)  $\leftarrow$  Tutor 2
- .....
- (Sam, 25)  $\leftarrow$  Student 1
- (Ethan, 26)  $\leftarrow$  Student 2
- .....

0	1	2	3	...	25	26	....
X	Marina	Kalkin	Cassidy		Elif	Grace	

# Running time for Direct Hashing

```
Class StudentDataBase {  
    Student[] allStudents = new Student [4294967268];  
    void add (Student s) {  
        int index = s.studentId;  
        allStudents[index]=s;  
    }  
    Student get (Student s) {  
        int index = s.studentId;  
        return allStudent[index];  
    }  
    void remove(Student s) {  
        int index = s.studentId;  
        allStudents[index] = null;  
    }  
}
```


# Space complexity is horrible :(

- Allocating 4 billion entries for a UCSD student database is extremely wasteful.
  - The universe of keys is far larger than the number we expect to ever want to store.
- What if we decrease the size of the table from 4 billion entries to storing, say, just 100,000?
- 100,000 is still far higher than the actual number of UCSD students, so there should be plenty of space.



# Collision

- With an array of only 100,000 entries, each student ID would **no longer have its own unique array index** -- multiple student IDs would have to “share” an index.
  - We call the “sharing” of an array index by 2 (or more) student IDs a **collision**.
- Whenever a collision occurs, we have to store the Student object “somewhere else” (more later).



# Collision

- However, if we're **clever** about **how** we assign array indices to student IDs, then collisions will **rarely** occur.
- We can still achieve  $O(1+..)$  add/find/remove time in the *average case*.



# Idea: Hash Tables

- A hash table consists of a large array of  $M$  “slots” (or “buckets”) to store the user’s data.
- A hash table also requires:
  1. Some way of converting from an object’s *key* into an *index* that specifies where that object should be stored. This is called a **hash function**.
  2. A method of handling **collisions**.

In order to ensure good performance,  $M$  (number of cells) must be bigger than  $N$ , the number of data the user will want to store.

# Do we have a collision?

M=5; (size of the hash table)

Student1 id is 1;

Student2 id is 4;

Student3 id is 8;

A: Yes

B: No

C: May be

# Hash function

- A *hash function* maps an object's *key* into an array *index*, i.e., a number from  $0 \dots M-1$ , where  $M$  is the number of entries in the hash table.
- **Example:**

```
int hashFucntion (int studentID){  
    return studentID % M; }
```

The modulus operator % divides studentID by M, and then returns the remainder. Examples:

- $3 \% 10 = ?$
- $107 \% 10 = 7$
- $7 \% 4 = 3$
- $16 \% 5 = 1$

# Do we have a collision?

$M=5$ ; (size of the hash table)

Student1 id is 1;

Student2 id is 4;

Student3 id is 8;

A: Yes

B: No

C: May be

# Hash function

- To be useful, a hash function must be **fast**
  - Its performance should not depend on the particular key.
  - $O(1+..)$  (otherwise how can we achieve desired  $O(1+..)$  performance? )

# Hash function

- To be useful, a hash function must be **fast**
  - Its performance should not depend on the particular key.
  - $O(1+..)$  (otherwise how can we achieve desired  $O(1+..)$  performance? )
- A hash function must also be **deterministic**:
  - Given the same key, it must always return the same array index.  
(Otherwise, how would we find something we stored earlier?)



# Hash function

- To be useful, a hash function must be **fast**
  - Its performance should not depend on the particular key.
  - $O(1+..)$  (otherwise how can we achieve desired  $O(1+..)$  performance? )
- A hash function must also be **deterministic**:
  - Given the same key, it must always return the same array index.  
(Otherwise, how would we find something we stored earlier?)
- A “good” hash function should also be **uniform**:
  - Each “slot”  $i$  in the array should be equally likely to be chosen as any other slot  $j$ .

# Is it a good hash function?

```
int hashFucntion (int studentID) {  
    return M/2; } //M is a size of a Hash Table
```

A: Yes, I like it

B: No, it is not fast

C: No, it is not deterministic

D: No, it is not uniform

# Hash function

- For instance, if  $M$  is 100000, then  $\text{studentID} \% 100000$  is simply the *last 5 digits* of the student ID, e.g.:

- `student1` with Student ID 0000013012 would map to index 13012.

- `student2` with Student ID 1234567890 would map to index 67890.

- These *indices* specify *where* in the array the students are *stored*.

Key (student ID)	Value (reference to Student object)
...	
13011	
13012	<code>student1</code>
13013	
...	
67889	
67890	<code>student2</code>
67891	
67892	
...	...

# Handling collisions

- Unfortunately, on occasion, there would be two (or more) `Student` objects who are “hashed” (mapped) into the same array slot.

```
studentID1 = 2200012345;  
studentID2 = 1926112345;  
hashTable.add(studentID1, student1);  
hashTable.add(studentID2, student2);
```


- This is called a *collision* -- two different `Student` objects map into the *same array index*.
- How do we handle these collisions?

# Handling collisions

- There are two principal ways of handling collisions:
  1. **Chaining** (aka **separate chaining**) -- at each slot in the array, instead of storing only a single element, we store a linked list of elements.
  2. **Open addressing** -- if `student5` “hashes” to array index 123, and array index 123 is already occupied, then we look for “another” index at which to store `student5`, e.g., 124.
    - Different schemes for determining “another index”.

# Chaining

- Each slot in the array contains not an object itself, but rather a pointer to the *head* of a *linked list* of objects which all map to the same index.




# Chaining

- When looking for a particular object, we must:

1. *Hash* the key to obtain the *index*.
2. *Search* the list for the correct object.

- This will still be *fast* as long as the linked lists are *short* (more later).






# Assume a good hash function

Is this possible?

A: Yes

B: No

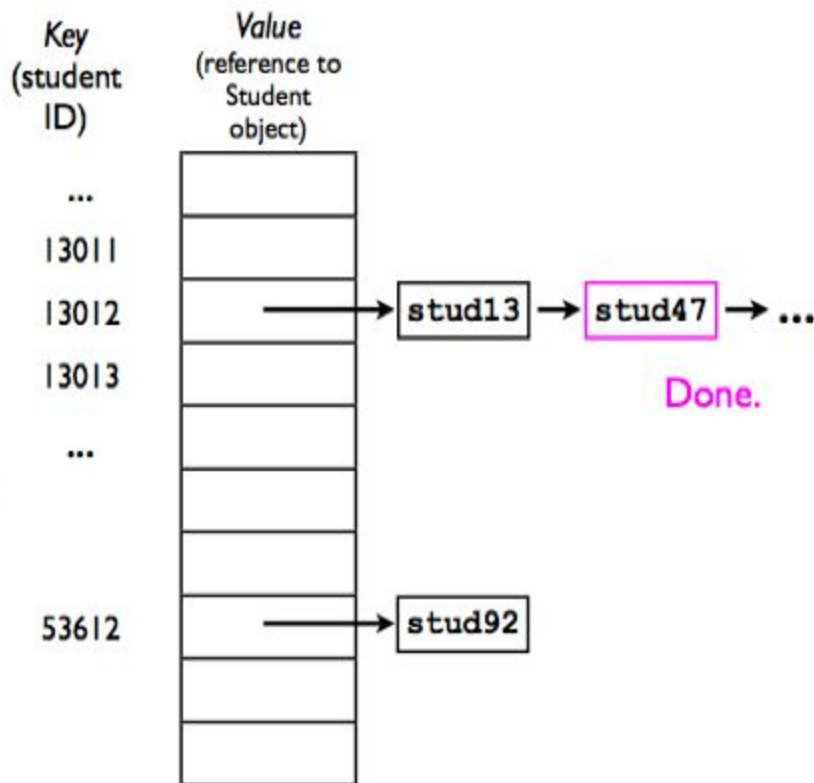




# Chaining

- For example, if we wish to find `stud47` with student ID 0925113012.

- Hash `stud47`'s student ID to determine the *index*.
- Jump to the head of the corresponding linked list.
- Traverse the linked list until we find `stud47`.



# Chaining: complexity

How fast are the **add/remove/find** operations for a hash table with chaining?

To find an object in a hash table, we must:

- **Hash the key.** A:  $O(1)$ , B:  $O(n)$ , C: unknown
- **Jump to that array index.** A:  $O(1)$ , B:  $O(n)$ , C: unknown
- **Traverse** through the linked list at that index (worst).

A:  $O(1)$

B:  $O(n)$

C: *Something else*

# Chaining: performance analysis

In the *worst case*, all  $N$  objects stored in the hash table will hash to the *same array index*

- This means that the linked list at that index will be  $N$  elements long.
- To find an arbitrary element in a linked list we need  $O(N)$  time.
- This is no better than just using a linked list by itself!

# Chaining: performance analysis

- However, in the average case, a hash table performs much better:
  - Given  $M$  slots in the array and  $N$  objects to store, the average list length for any array slot is  $N/M$ . (Goes without proof)
  - Then, the average time to access any arbitrary object is
  - $O(1+N/M)$ . (Goes without proof).
- Now, suppose that we always make sure that  $M > N$ :
  - Then  $N/M < 1$ .
  - Therefore, average-case time cost is  $O(1+N/M) = O(1)$



# Load factor

Load factor  $\alpha = N/M$  of the hash table.

**The proper way** to report a running time for hash tables is:

$O(1+\alpha)$ , where  $\alpha$  is a load factor. If  $\alpha$  becomes too large, then the running time is not a constant anymore.

For separate chaining load factor of 1 is OK. If it becomes larger:

**double the size of the table and rehash.**

# reminders

\*mic

PA8, 9

# Open addressing

# Idea

- In a case of a collision, we look for another available cell.
- Must be **deterministic**
  - If we want to find an element, we should visit the **same** cells.



# Simplest: linear probing

- If **hashFunction(key)** maps into an index  $i$  that is already occupied, then try  $i+1$ .
- If that doesn't work, try  $i+2$ ,  $i+3$ , ..., etc.
- If we get to  $M-1$ , we want to “wrap around” back to 0.
- The index of the  $j$ th **probe** (where  $j$  starts at 0) is given by the expression  $(i+j) \% M$
- You can think of the *probe* as one step in your search

# Exercise

- The index of the  $j$ th probe (where  $j$  starts at 0) is given by the expression  $(i+j) \% M$
- Insert the given numbers in the hash table ( $M=7$ ), resolving collisions by linear probing ( use  $\%M$  for hash function)

1, 15, 2, 7, 14, 22

$\% 7$

1, 15, 2, 7, 14, 22

Click to add text

→ 0	7
→ 1	1
→ 2	15
→ 3	2
→ 4	14
→ 5	22

$$1 \% 7 \rightarrow 1$$

$$15 \% 7 \rightarrow 1$$

$$2 \% 7 \rightarrow 2$$

$$7 \% 7 \rightarrow 0$$

$$14 \% 7 \rightarrow 0$$

$$22 \% 7 \rightarrow 1$$

# Primary clustering

- If too many keys hash to the same index -- *or to nearby indices* -- then the linear probing may become expensive.
- Consider the hash table to the right:
  - 13011-13016 are already occupied.
  - If we want to add another student `student7` who also hashes to 13011, then we have to step through 7 elements.
  - The longer the cluster, the higher the time cost for add/find/remove.

Key (student ID)	Value (reference to Student object)
...	
13011	<code>student5</code>
13012	<code>student1</code>
13013	<code>student9</code>
13014	<code>student8</code>
13015	<code>student3</code>
13016	<code>student4</code>
13017	
...	

# Removing an element

Suppose we want to remove student1 from the hash table:

```
Table[13011] = null;
```

**Good idea?**

A: Yes.

B: No

C: Seems like a bad idea but I do not know why.

Key (student ID)	Value (reference to Student object)
...	
13011	student1
13012	student2
13013	
...	
53612	

# Removing an element

- Suppose we remove `student1` from the hash table.
- If we later search for `student2`, we will still hash to 13011, but find that it is *empty*.
- Does that mean `student2` is not contained in the hash table?
- No -- but we have to *record* that somehow.

Key (student ID)	Value (reference to Student object)
...	
13011	
13012	<code>student2</code>
13013	
...	
53612	

# Removing an element (lazy deletion)

- One method of recording that an element was deleted is a *bridge*, a special element that indicates “empty, but keep looking.”
- If we later add another element, say `student5` that hashes to `13011`, then we can *replace* the bridge with a real `Student` object.

Key (student ID)	Value (reference to Student object)
...	
13011	(bridge)
13012	student2
13013	
...	
53612	

# Removing an element (lazy deletion)

- One method of recording that an element was deleted is a *bridge*, a special element that indicates “empty, but keep looking.”
- If we later add another element, say `student5` that hashes to `13011`, then we can *replace* the bridge with a real Student object.

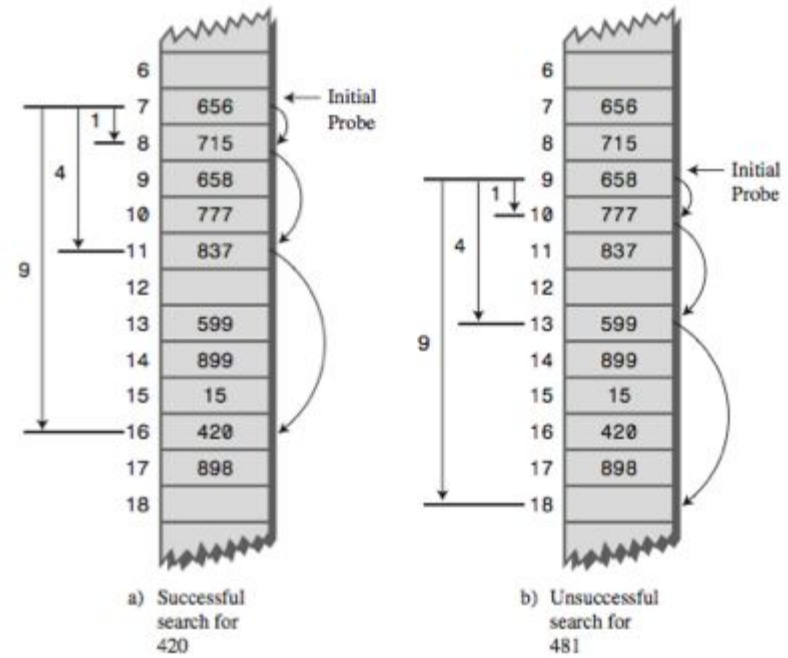
Key (student ID)	Value (reference to Student object)
...	
13011	<code>student5</code>
13012	<code>student2</code>
13013	
...	
53612	



# Quadratic probing

- Quadratic probing is an attempt to keep clusters from forming. The idea is to probe more widely separated cells, instead of those adjacent to the primary hash site.

- In quadratic probing, probes go to  $x+1$ ,  $x+4$ ,  $x+9$ ,  $x+16$ ,  $x+25$ , and so on. The distance from the initial probe is the square of the step number:  $x+1^2$ ,  $x+2^2$ ,  $x+3^2$ ,  $x+4^2$ ,  $x+5^2$ , and so on



# The Problem with Quadratic Probes

- Quadratic probes suffer from a different and more subtle clustering problem.
- Let's say 184, 302, 420, and 544 all hash to 1 and are inserted in this order.
- What is the hash table after we insert these numbers? (assume the table is large enough)

- 184, 302, 420, and 544

# Secondary clustering

- Quadratic probes suffer from a different and more subtle clustering problem.
  - This occurs because all the keys that hash to a particular cell follow the same sequence in trying to find a vacant space.
- Let's say 184, 302, 420, and 544 all hash to 1 and are inserted in this order. Each additional item with a key that hashes to 1 will require a longer probe.
  - This phenomenon is called secondary clustering.
- Secondary clustering is not a serious problem, but quadratic probing is not often used because there's a slightly better solution.

# Double Hashing

# Double Hashing

**Idea:** generate probe sequences that depend on the **key** instead of being the same for every key.

Then numbers with *different* keys that hash to the same index will use *different* probe sequences.

The solution is to hash the key a **second** time, using a different hash function, and use the result as the step size.

- $h(k, i) = (h_1(k) + i h_2(k)) \bmod m.$ 
  - $h_1$  and  $h_2$  are auxiliary hash functions.
  - $i$  is the probe step

**Example:  $h(k, i) = (h_1(k) + i h_2(k)) \bmod m$ .**

- Assume  $M = 13$  (size of the table)
- $h_1(k) = k \bmod 13$
- $h_2(k) = 1 + (k \bmod 11)$ .

Let's insert 79, 69, 98, 72, 14, 50 using double hashing.

$h(k, i) = (h_1(k) + i h_2(k)) \bmod m;$  79, 69, 98, 72, 14, 50;  $h_1(k) = k \bmod 13$ ,  $h_2(k) = 1 + (k \bmod 11)$



# Restrictions on $h_2$

- **Experience** has shown that this secondary hash function must have certain characteristics:
  1. It must not be the same as the primary hash function.
  2. It must never output a 0 (otherwise, there would be no step; every probe would land on the same cell, and the algorithm would go into an endless loop).

# Analysis of open-address hashing

- Load factor  $\alpha = N/M$  of the hash table. Since at most one element occupies each slot and  $M \geq N$ ,  $\alpha \leq 1$ .
- Facts (without proof).
  - For linear probing the load factor can't be more than  $2/3$  (under  $1/2$  is the best)
  - For quadratic probing and double hashing load factor of  $2/3$  is OK.
  - If the load factor increases, you must re-hash: create a larger array and **re-hash** every element into a new hash table.