

# DSC 40B

## *Theoretical Foundations II*


Lecture 17 | Part 1

### Kruskal's Algorithm

# Last Time: Minimum Spanning Tree

- ▶ The **minimum spanning tree** problem is as follows:
  - ▶ GIVEN: A weighted, undirected graph  $G = (V, E, \omega)$ .
  - ▶ COMPUTE: a spanning tree of  $G$  with minimum cost (i.e., minimum total edge weight).

# Example




# Last Time: Building MSTs

- ▶ How do we build a MST efficiently?
- ▶ We'll adopt a **greedy** approach.
  - ▶ Build a tree edge-by-edge.
  - ▶ At every step, doing what looks best at the moment.
- ▶ This strategy isn't guaranteed to work in all of life's situations, but it works for building MSTs.

# Two Greedy Approaches


- ▶ We'll look at two greedy algorithms:
  - ▶ Last Time: Prim's Algorithm
  - ▶ Today: Kruskal's Algorithm
- ▶ Differ in the order in which edges are added to tree.
- ▶ Also differ in time complexity.

# Prim's Algorithm, Informally



- ▶ Start by picking any node to add to “tree”,  $T$ .
- ▶ While  $T$  is not a spanning tree, greedily add **lightest** edge from a node in  $T$  to a node not in  $T$ .
  - ▶ “lightest” = edge of smallest weight

# Kruskal's Algorithm, Informally



- ▶ Start with empty forest:  $T = (V, E_{\text{mst}})$ , where  $E_{\text{mst}} = \emptyset$ .
- ▶ Loop through edges in increasing order of weight.
  - ▶ If edge does not create a cycle in  $T$ , add it to  $T$ .
  - ▶ If  $T$  is a spanning tree, break.

# Being Greedy

- ▶ Prim: add the **node** with smallest estimated cost and update neighbors.
  - ▶ Works locally, “grows” a connected tree.
- ▶ Kruskal: add the **edge** with smallest weight.
  - ▶ As long as it doesn't make a cycle.
  - ▶ Edge can be anywhere in graph.



# Kruskal's Algorithm (Pseudocode)

```
def kruskal(graph, weights):  
    mst = UndirectedGraph()  
  
    # sort edges in ascending order by weight  
    sorted_edges = sorted(graph.edges, key=weights)  
  
    for (u, v) in sorted_edges:  
        # if u and v are not already connected  
        if ...:  
            mst.add_edge(u, v)  
  
            # (optional) if mst is now a spanning tree, break  
            if len(mst.edges) == len(graph.nodes) - 1:  
                break  
  
    return mst
```

# Checking for Connectivity

- ▶ Each iteration: check if  $u$  and  $v$  are connected in  $T = (V, E_{\text{mst}})$ .
- ▶ We *could* do a DFS/BFS on each iteration...
  - ▶  $\Theta(V + E_{\text{mst}}) = \Theta(V)$  each time.
  - ▶ **Expensive!**
- ▶ Remember:
  - ▶ If you're computing something once, use a fast algorithm.
  - ▶ If you're computing it repeatedly, consider a **data structure**.

# Disjoint Set Forests

- ▶ Represent a collection of disjoint sets.

$\{\{1, 5, 6\}, \{2, 3\}, \{0\}, \{4\}\}$

- ▶ `.union(x, y)`: Union the sets containing `x` and `y`.
- ▶ `.in_same_set(x, y)`: Return **True/False** if `x` and `y` are in the same set.<sup>1</sup>

---

<sup>1</sup>Usually implemented as a `.find(x)` method returning representative of set containing `x`.


$\{0, 3, 1, 4\}, \{2, 5\}$

## Example

```
»» # create a DSF with {{0}, {1}, {2}, {3}, {4}, {5}}
»» dsf = DisjointSetForest([0, 1, 2, 3, 4, 5])
»» dsf.union(0, 3)
»» dsf.union(1, 4)
»» dsf.union(3, 1)
»» dsf.union(2, 5)
»» # dsf now represents {{0, 1, 3, 4}, {2, 5}}
»» dsf.in_same_set(0, 3)
True
»» dsf.in_same_set(0, 2)
False
```

# Disjoint Set Forests

- ▶ Operations take  $\Theta(\alpha(n))$  time, where  $n$  is number of objects in collection.
- ▶  $\alpha(n)$  is the **inverse Ackermann function**.
- ▶ It grows very, **very** slowly.
- ▶ Essentially constant time.



# Disjoint Set Forests

- ▶ Can be used to keep track of CCs of a **dynamic graph**.
- ▶ Nodes of CCs are disjoint sets.
  - ▶ Add an edge  $(u, v)$ : `.union(u, v)`
  - ▶ Check if  $u$  and  $v$  are connected:  
`.in_same_set(u, v)`
- ▶ To check if  $u, v$  are already connected:
  - ▶ BFS/DFS:  $\Theta(V)$  each time.
  - ▶ DSF:  $\Theta(\alpha(V))$  each time (essentially  $\Theta(1)$ ).

# Kruskal's Algorithm

```
def kruskal(graph, weights):  
    mst = UndirectedGraph()  
  
    # place each node in its own disjoint set  
    components = DisjointSetForest(graph.nodes)  
  
    # sort edges in ascending order by weight  
    sorted_edges = sorted(graph.edges, key=weights)  
  
    for (u, v) in sorted_edges:  
        if not components.in_same_set(u, v):  
            mst.add_edge(u, v)  
            components.union(u, v)  
  
        # (optional) if mst is now a spanning tree, break  
        if len(mst.edges) == len(graph.nodes) - 1:  
            break  
  
    return mst
```

# Time Complexity

```
def kruskal(graph, weights):  
    mst = UndirectedGraph()
```

```
    # place each node in its own disjoint set  
    components = DisjointSetForest(graph.nodes)
```

$\{ \Theta(V)$

```
    # sort edges in ascending order by weight  
    sorted_edges = sorted(graph.edges, key=weights)
```

$\{ \Theta(E \log E)$

```
    for (u, v) in sorted_edges:  
        if not components.in_same_set(u, v):  
            mst.add_edge(u, v)  
            components.union(u, v)
```

$\left. \begin{array}{l} \text{loop iterates } \leq |E| \text{ times} \\ \text{time per iter: } \alpha(V) \end{array} \right\} \text{total } \alpha(|E| \alpha(V))$

```
    # (optional) if mst is now a spanning tree, break  
    if len(mst.edges) == len(graph.nodes) - 1:  
        break
```

```
    return mst
```

$\Theta(E \log E + V)$



# Time Complexity

- ▶ Assume graph is connected. Then  $E = \Omega(V)$ .
- ▶ Kruskal's takes  $\Theta(E \log E) = \Theta(E \log V)$  time.
  - ▶ Dominated by sorting the edges.
- ▶ Note: if graph disconnected, Kruskal's produces a **minimum spanning forest**.

# DSC 40B

## *Theoretical Foundations II*

Lecture 17 | Part 2

**Kruskal v. Prim**

# Kruskal v. Prim

- ▶ Both algorithms for computing MSTs.
- ▶ Which is “better”?
- ▶ There's no clear winner.

# Time Complexity

- ▶ Prim:
  - ▶ Binary heap:  $\Theta(V \log V + E \log V)$
  - ▶ Fibonacci heap:  $\Theta(V \log V + E)$
- ▶ Kruskal:  $\Theta(E \log V)$
- ▶ If the graph is dense,  $E = \Theta(V^2)$ , and Prim's with Fibonacci heap "wins".
  - ▶  $\Theta(V^2)$  versus  $\Theta(V^2 \log V)$ .

## Not so fast...

- ▶ Fibonacci heaps are hard to implement, high overhead.
- ▶ Prim's will be faster for very large dense graphs.
- ▶ But Kruskal's may be faster for smaller dense graphs.
- ▶ The right choice depends on your application.

## Main Idea

Asymptotic time complexity isn't everything. For small inputs, the "inefficient" algorithm may beat the "efficient" one. There's also ease of implementation to consider.

# DSC 40B


## *Theoretical Foundations II*

Lecture 17 | Part 3

### **MSTs and Clustering**

# Clustering


Goal: identify the groups in data. Example:





# Clustering, Formalized


- ▶ We frame as an optimization problem.
  - ▶ GIVEN:  $n$  data points.
  - ▶ GOAL: assign color to each point (red or blue) to maximize the distance between the closest pair of red and blue points.



**Bad Clustering**

# Clustering, Formalized

- ▶ We frame as an optimization problem.
  - ▶ GIVEN:  $n$  data points.
  - ▶ GOAL: assign color to each point (red or blue) to maximize the distance between the closest pair of red and blue points.





**Good Clustering**

# Brute Force Solution

- ▶ Try all possible assignments; return best.
- ▶ If there are  $n$  data points, there are  $\Theta(2^n)$  assignments.
- ▶ Exponential time; very slow. Practical only for  $\sim 50$  data points.
- ▶ Instead, we will turn it into a graph problem.

# Distance Graphs

- ▶ Given  $n$  data points,  $p_1, p_2, \dots, p_n$ , create complete graph with  $V = \{p_1, \dots, p_n\}$ .
- ▶ Set weight of edge  $(p_i, p_j) = \text{dist}(p_i, p_j)$ .
- ▶ The result is a weighted, undirected **distance graph**.




## Main Idea

We can always think of a set of points in a (metric) space as a weighted distance graph. This is a **very** important idea, because it allows us to use our graph algorithms!

# Clustering with MSTs

- ▶ Given  $n$  data points and a number of clusters,  $k$ :
  - ▶ Create distance graph  $G$ .
  - ▶ Run Kruskal's Algorithm on  $G$  until there are only  $k$  components.



- ▶ The resulting connected components are the **clusters**.
- ▶ This is known as **single-linkage clustering**.

# Example

# Single-Linkage Clustering

- ▶ Time complexity of single-linkage is determined by Kruskal's Algorithm:  $\Theta(E \log E)$ .
- ▶ Since distance graph is complete,  $E = \Theta(V^2)$ , and so
$$\Theta(E \log E) = \Theta(V^2 \log V) = \Theta(n^2 \log n)$$
- ▶ Practically, can cluster ~ 10,000 points.



# Summary

- ▶ We started the quarter with a brute force solution.
  - ▶ Took  $\Theta(2^n)$  time, only feasible for a few dozen points.
- ▶ We've now reframed the problem using graph theory.
  - ▶ Now only  $\Theta(n^2 \log n)$  time!
  - ▶ Feasible for tens of thousands of points.

# Why Algorithms?

- ▶ Data scientists use computers as tools.
- ▶ But solving a problem isn't just about coding it up.
- ▶ You need to know how to analyze your code and use the right algorithms and data structures to make your solution efficient.